

**AD CLICK-THROUGH RATE PREDICTION FOR DIGITAL
MARKETING**

INTERNSHIP PROJECT REPORT



Submitted by

**AALIA FATHIMA W
VC0N100565138999
DATA SCIENCE INTERN**

**BACHELOR OF ENGINEERING
IN
COMPUTER SCIENCE AND ENGINEERING**

ABSTRACT

In the digital era, online advertising is vital for businesses, but accurately predicting the likelihood of a user clicking on an advertisement (CTR) remains a key challenge. Ineffective targeting wastes ad budgets and reduces user engagement. This project aims to build a CTR Prediction Model using machine learning to predict the probability of a user clicking on an ad by leveraging historical user behavior, demographics, and ad-related features. The workflow includes data collection, preprocessing, feature engineering, model training, evaluation, and deployment. Various machine learning algorithms, including Logistic Regression, Random Forest, and Gradient Boosting, are applied and compared using evaluation metrics such as MAE, RMSE, and accuracy. The final model is deployed using the Pickle module for integration with web applications or ad-serving systems for real-time predictions. This project benefits advertisers and ad-tech platforms by enabling personalized ad delivery, improving user engagement, and optimizing ad spend through accurate CTR prediction. It demonstrates how machine learning can effectively address real-world challenges in digital marketing, showcasing the potential of data-driven strategies to improve ad targeting, user satisfaction, and business outcomes in the digital advertising ecosystem.

TABLE OF CONTENTS

CHAPTER NO.	TITLE	PAGE NO.
	ABSTRACT	i
	LIST OF FIGURES	iii
	LIST OF ABBREVIATIONS	iv
1.	INTRODUCTION	01
	1.1 Overview	01
	1.2 Objective	01
	1.3 Scope	02
	1.4 Problem Statement	03
2.	LITERATURE REVIEW	04
3.	SYSTEM ANALYSIS	08
	3.1 Existing System	08
	3.2 Proposed System	09
	3.3 Limitations of the Project	10
	3.4 Development Environment	11
4.	SYSTEM DESIGN	12
	4.1 Introduction	12
	4.2 Data Flow Diagram (DFD)	12
	4.3 System Architecture	16
	4.4 Libraries and Tools	22

CHAPTER NO.	TITLE	PAGE NO.
	4.5 Modules	25
	4.6 Pickle/Joblib Module	29
5.	CONCLUSION	31
	5.1 Future Scope	31
	5.2 Conclusion	32
6.	APPENDICES	34
	6.1 Output	34
7.	REFERENCES	36

LIST OF FIGURES

FIGURE NO.	FIGURE DESCRIPTION	PAGE NO.
4.2.1	Level 0 DFD	13
4.2.2	Level 1 DFD	16
4.3	Architecture Diagram	17
6.1.1	Data Input 1	34
6.1.2	Data Input 2	34
6.1.3	Poor Targeting	35
6.1.4	Excellent Targeting	35
6.1.5	Moderate Targeting	35

LIST OF ABBREVIATIONS

CTR – Click-Through Rate

ML – Machine Learning

CSV – Comma Separated Values

GUI – Graphical User Interface

API – Application Programming Interface

MAE – Mean Absolute Error

RMSE – Root Mean Square Error

CHAPTER 1

INTRODUCTION

1.1 OVERVIEW

In today's data-driven digital landscape, online advertising has become a cornerstone for businesses aiming to reach targeted audiences and maximize their return on investment. Click-Through Rate (CTR), which measures the ratio of users who click on an advertisement to the total number of users who view it, is a critical metric in evaluating the effectiveness of online ad campaigns. However, accurately predicting whether a user will click on an advertisement remains a significant challenge due to the complexity of user behavior, diverse demographics, and varying advertisement contexts. The CTR Prediction Model developed in this project utilizes machine learning techniques to predict the likelihood of a user clicking on an advertisement based on their historical behavior, demographic information, and advertisement features. This data science project leverages a systematic approach, involving data collection, preprocessing, feature engineering, model training, and evaluation to build a robust and scalable prediction system capable of handling real-world advertising datasets. The model is designed to learn from historical click patterns and contextual information, thereby enabling more accurate and personalized ad targeting, optimizing ad spend, and improving the user experience by delivering relevant advertisements.

1.2 OBJECTIVE

The primary objective of this project is to apply machine learning algorithms to predict the Click-Through Rate (CTR) for online advertisements, thereby enhancing the accuracy and efficiency of digital marketing campaigns. This project, titled “ Ad Click-Through Rate Prediction for Digital Marketing,” aims to develop a predictive model capable of analyzing user behavior, demographic data, and advertisement attributes to estimate the probability of a user clicking on a specific advertisement. By leveraging data science and machine learning methodologies, the CTR Prediction Model enables advertisers

and digital marketing professionals to make informed decisions regarding ad placements, budget allocations, and campaign strategies. This project also serves as a practical implementation of theoretical knowledge acquired in data science, machine learning, and predictive analytics, providing hands-on experience in addressing real-world challenges faced by the digital advertising industry.

1.3 SCOPE

The scope of the CTR Prediction Model project encompasses the end-to-end development of a machine learning system designed to predict the likelihood of user clicks on online advertisements. This includes the design and implementation of data pipelines for collecting and preprocessing user behavior and demographic data, along with advertisement-specific features. The project involves exploratory data analysis, feature engineering to enhance model performance, training multiple machine learning algorithms, and evaluating their effectiveness using standard metrics such as accuracy, precision, recall, and error rates. Additionally, the project focuses on deploying the trained model using serialization techniques like Pickle to enable seamless integration into web applications and ad-serving platforms for real-time predictions. While the system does not replace the need for continuous marketing strategy refinement, it serves as a valuable tool for advertisers and marketing agencies seeking to optimize their campaigns through data-driven insights. Future enhancements to the project may include the integration of advanced deep learning models, real-time data streaming, and adaptive learning mechanisms to further improve prediction accuracy. The CTR Prediction Model demonstrates the impactful role of machine learning in digital marketing and showcases how technology can be leveraged to create intelligent systems that benefit businesses while enhancing the user experience through personalized and relevant advertisements.

1.4 PROBLEM STATEMENT

In the dynamic landscape of digital marketing, businesses invest substantial resources in online advertising with the goal of reaching potential customers and maximizing engagement. However, one of the major challenges faced by advertisers and marketing agencies is accurately predicting whether a user will click on a given advertisement. Existing ad-serving systems often rely on simplistic rule-based targeting or generic user segmentation, which fails to capture the complex patterns underlying user behavior, demographic factors, and the contextual attributes of advertisements.

This lack of precision in targeting leads to several challenges, including inefficient budget utilization, reduced campaign effectiveness, and poor user experience due to irrelevant advertisements. Advertisers may end up displaying ads to users with low interest or likelihood of engagement, resulting in wasted impressions and lower return on investment. Additionally, users often experience ad fatigue when exposed to irrelevant advertisements repeatedly, reducing overall trust in online advertising platforms.

The problem addressed in this project is to develop a robust, machine learning-based Click-Through Rate (CTR) Prediction Model that can accurately predict the likelihood of a user clicking on an advertisement by analyzing historical user behavior, demographic data, and advertisement-specific features. By addressing this problem, the project aims to enhance the efficiency and effectiveness of digital advertising campaigns, enabling advertisers to make data-driven decisions for ad targeting, improve user engagement, and optimize advertising costs.

This problem statement forms the foundation of the project, emphasizing the need for an intelligent system that leverages data science and machine learning to bridge the gap between advertisers and potential customers in a scalable, accurate, and user-centric manner.

CHAPTER 2

LITERATURE REVIEW

1. **Title:** Predicting Click-Through Rate with Machine Learning for Online Advertising

Author: He et al.

Year: 2014

Goal: To improve CTR prediction accuracy in online advertising using machine learning models.

Algorithm: Logistic Regression, Gradient Boosting Decision Trees (GBDT)

Description: The study explored the use of machine learning algorithms to predict user clicks on advertisements, focusing on the scalability of models to handle billions of samples in real-time ad systems. Logistic Regression provided a strong baseline, while GBDT captured nonlinear feature interactions, improving predictive performance and ad targeting efficiency.

2. **Title:** Practical Lessons from Predicting Clicks on Ads at Facebook

Author: Gupta et al. (Facebook)

Year: 2014

Goal: To apply scalable learning systems for CTR prediction in a production environment.

Algorithm: Logistic Regression with Feature Engineering

Description: This work detailed the engineering practices adopted at Facebook for large-scale CTR prediction, emphasizing the significance of feature selection, handling categorical variables, and balancing data in training sets. It demonstrated how logistic regression, when combined with effective feature engineering, could achieve reliable CTR predictions while maintaining computational efficiency for real-time serving.

3. **Title:** Predicting Click-Through Rate: A Study of Ad Impressions

Author: Richardson et al. (Microsoft Research)

Year: 2007

Goal: To enhance ad selection by predicting CTR using user and ad features.

Algorithm: Logistic Regression, Decision Trees

Description: This foundational study investigated the effectiveness of machine learning techniques for CTR prediction using features such as ad text, user demographics, and query-ad relevance. It demonstrated that machine learning methods significantly outperformed rule-based approaches in predicting CTR, paving the way for their adoption in commercial ad-serving systems.

4. **Title:** Deep Learning for Predicting Click-Through Rates of Ads

Author: Zhang et al.

Year: 2016

Goal: To evaluate the potential of deep learning models in CTR prediction tasks.

Algorithm: Deep Neural Networks (DNN)

Description: The paper presented the use of deep neural networks to automatically learn high-order feature interactions without manual feature engineering for CTR prediction. The results showed that DNNs could achieve higher prediction accuracy compared to traditional models, demonstrating the effectiveness of deep learning in capturing complex patterns in large-scale advertising data.

5. **Title:** A Survey on Click-Through Rate Prediction for Online Advertising

Author: Juan et al.

Year: 2015

Goal: To summarize methods, challenges, and datasets used in CTR prediction.

Algorithm: Various (Logistic Regression, Factorization Machines, Neural Networks)

Description: This survey provided a comprehensive overview of machine learning methods applied in CTR prediction, discussing dataset challenges, evaluation metrics, and the importance of feature engineering. It highlighted how different algorithms are suited to different ad systems based on data volume and prediction latency requirements, offering guidance for designing efficient CTR prediction systems.

6. Title: Predicting Ad Click-Through Rates via Feature-based Collaborative Filtering

Author: Gai et al.

Year: 2017

Goal: To improve CTR prediction by combining feature-based and collaborative filtering approaches.

Algorithm: Feature-based Collaborative Filtering, Logistic Regression

Description: This study proposed using collaborative filtering methods alongside feature-based models to predict CTR, leveraging user and item latent features to enhance the predictive power. By incorporating user-ad interaction data, the model was able to improve personalization and increase CTR prediction accuracy over traditional methods.

7. Title: DeepFM: A Factorization-Machine based Neural Network for CTR Prediction

Author: Guo et al.

Year: 2017

Goal: To integrate factorization machines with deep neural networks for CTR prediction without manual feature engineering.

Algorithm: DeepFM (Deep Neural Networks + Factorization Machines)

Description: The paper introduced DeepFM, which combines the strengths of factorization machines in modeling low-order feature interactions and deep neural networks for high-order interactions. This hybrid approach allowed the model to automatically learn feature interactions efficiently, resulting in higher CTR prediction performance on benchmark datasets.

8. Title: Predicting Click-Through Rates of New Ads Using Collaborative Filtering Models

Author: McMahan et al.

Year: 2013

Goal: To address the cold-start problem in CTR prediction for new ads.

Algorithm: Collaborative Filtering, Logistic Regression

Description: This research tackled the challenge of predicting CTR for newly introduced advertisements with limited historical data by utilizing collaborative filtering techniques to leverage user similarity and ad content similarity. The approach enabled the system to provide reasonable CTR predictions even for ads without prior click history, improving ad serving in dynamic environments.

9. **Title:** Field-aware Factorization Machines for CTR Prediction

Author: Juan et al.

Year: 2016

Goal: To enhance CTR prediction using field-aware factorization machines to model field interactions efficiently.

Algorithm: Field-aware Factorization Machines (FFM)

Description: The study proposed FFM, which extends traditional factorization machines by allowing each feature to interact differently with other features depending on their fields (e.g., user field, ad field). This improved the ability to capture inter-field feature interactions crucial in CTR prediction tasks, leading to significant improvements in ad click prediction accuracy.

10. **Title:** Predicting Click-Through Rate with Factorization Machines in Display Advertising

Author: Rendle et al.

Year: 2012

Goal: To apply factorization machines to CTR prediction in large-scale display advertising datasets.

Algorithm: Factorization Machines (FM)

Description: The paper demonstrated the use of factorization machines for CTR prediction, emphasizing their capability to model interactions between sparse and high-dimensional categorical features commonly found in advertising datasets. FM models outperformed traditional logistic regression in experiments on real-world advertising data, highlighting their effectiveness in handling feature sparsity and improving CTR prediction.

CHAPTER 3

SYSTEM ANALYSIS

3.1 EXISTING SYSTEM

In the current digital marketing landscape, most online advertising platforms utilize rule-based or heuristic approaches for displaying advertisements to users. These existing systems typically rely on broad targeting parameters such as user location, general interests, age group, and time of ad display without deeply analyzing user-specific behavior patterns and ad context. As a result, the targeting is often generic, leading to low click-through rates, inefficient ad spending, and reduced campaign effectiveness.

Some systems use basic statistical methods or simplistic logistic regression models without comprehensive feature engineering or advanced algorithmic tuning. These models often fail to capture the complex, nonlinear relationships between user demographics, behavioral history, and ad-specific attributes, which are critical for accurately predicting the likelihood of a user clicking on an advertisement. Additionally, these systems struggle to handle high-dimensional and sparse categorical data effectively, limiting their predictive capabilities in large-scale advertising environments.

The limitations of the existing system can be summarized as:

- Inability to personalize ad delivery effectively for individual users.
- Limited feature extraction and interaction modeling, leading to lower prediction accuracy.
- Static targeting strategies that do not adapt to dynamic user behavior.
- Wasted ad impressions and budget due to irrelevant ad displays.
- Poor user experience due to repetitive and non-contextual advertisements.

These challenges highlight the need for an advanced, machine learning-based CTR Prediction Model capable of learning complex user-ad interaction patterns and providing accurate predictions to optimize ad targeting, which the proposed system aims to address.

3.2 PROPOSED SYSTEM

The proposed system aims to build a machine learning-based Click-Through Rate (CTR) Prediction Model that accurately predicts the likelihood of a user clicking on an advertisement using user behavior data, demographic attributes, and advertisement-specific features.

The system architecture consists of the following major components:

- **Data Collection:** Importing datasets containing user behavior, demographic data, and ad attributes.
- **Data Preprocessing:** Handling missing values, outlier detection, and feature encoding.
- **Exploratory Data Analysis (EDA):** Visualizing data distributions and understanding feature correlations.
- **Model Building:** Training machine learning models including Logistic Regression, Random Forest, and XGBoost.
- **Model Evaluation:** Measuring model performance using metrics such as Accuracy, MAE, and RMSE.
- **Deployment:** Saving the final model using Pickle for real-time inference in production.

By implementing this systematic pipeline, the proposed system enables advertisers to predict CTR with higher accuracy, ensuring that ads are displayed to users who are most likely to engage with them. This enhances user experience by delivering relevant ads while maximizing the return on investment for advertisers.

3.3 LIMITATIONS OF THE PROJECT

While the CTR Prediction Model using Machine Learning significantly improves the accuracy of ad click predictions and optimizes ad targeting, it also has certain limitations:

1. **Data Dependency:** The accuracy of the model heavily depends on the quality and quantity of historical data available. Insufficient or biased data can lead to inaccurate predictions.
2. **Cold Start Problem:** For new users or new advertisements with no prior interaction history, the model may struggle to predict CTR accurately due to a lack of training data related to those entities.
3. **Feature Engineering Challenges:** The effectiveness of the model relies on appropriate feature selection and engineering, which can be complex and time-consuming when dealing with high-dimensional and sparse categorical data.
4. **Model Interpretability:** Advanced models like Random Forest or XGBoost, while accurate, may function as black-box models, making it difficult to interpret why a particular prediction was made.
5. **Dynamic User Behavior:** User interests and behaviors can change over time, requiring periodic retraining of the model to maintain accuracy, which may increase maintenance overhead.
6. **Computational Resource Requirements:** Handling large advertising datasets and training complex models can require substantial computational resources and time, especially during hyperparameter tuning.
7. **Real-Time Prediction Latency:** Integrating the model into live ad-serving systems for real-time prediction may introduce latency if not optimized properly.
8. **Generalization Issues:** The model trained on specific datasets may not generalize well across different advertising platforms or regions without proper retraining and adaptation.

These limitations indicate the areas where further research, optimization, and system enhancements can be implemented to improve the efficiency, scalability, and effectiveness of the CTR Prediction Model in practical deployments.

3.4 DEVELOPMENT ENVIRONMENT

The CTR Prediction Model is developed in a Python-based data science environment, ensuring scalability, reproducibility, and efficient experimentation.

Programming Language:

- Python

Libraries and Frameworks:

- Pandas, NumPy: For data manipulation and numerical operations.
- Matplotlib, Seaborn: For data visualization.
- Scikit-learn: For machine learning model building, evaluation, and preprocessing.
- Pickle: For model deployment.

Development Tools:

- Kaggle Notebook / VS Code for development
- Git for version control.

System Requirements:

- RAM: 8GB minimum (16GB recommended for large datasets)
- Processor: Intel i5 or higher
- OS: Windows, macOS, or Linux

Dataset Source:

- Publicly available advertising datasets in kaggle.

The development environment ensures seamless integration of various modules in the project, enabling efficient experimentation and iterative improvements in model performance while maintaining organized, readable, and scalable code for future enhancements.

CHAPTER 4

SYSTEM DESIGN

4.1 INTRODUCTION

The CTR Prediction Model System is designed to accurately predict whether a user will click on an advertisement based on their behavior, demographics, and advertisement features. This system combines machine learning techniques with a user-friendly web interface to provide real-time, data-driven insights for optimizing digital marketing strategies. The system architecture is modular, consisting of data collection, preprocessing, feature engineering, model training, evaluation, and deployment components. Using libraries like pandas, scikit-learn, and Streamlit, the system ensures seamless workflow from raw data ingestion to user-facing predictions. It is designed for scalability, enabling handling of large advertising datasets while maintaining accuracy and efficiency in predictions.

4.2 DATA FLOW DIAGRAM (DFD)

The Data Flow Diagram (DFD) visually represents the logical flow of data in the CTR Prediction Model system, illustrating how raw user and ad data are transformed into actionable CTR predictions. It helps in understanding the system workflow clearly and systematically, aiding in implementation, debugging, and future maintenance. The DFD is structured into two levels:

4.2.1 Level 0 DFD (Context Diagram)

The Level 0 DFD provides a high-level overview of the system as a single process and shows the data flow between the system and external entities.

External Entities:

- **User Data Source:** Provides user demographics and behavior data.
- **Ad Data Source:** Provides advertisement details (topic, region).
- **User (via Streamlit App):** Inputs new data to predict CTR.

Process:

- CTR Prediction System: The main processing unit that takes user/ad data, preprocesses it, applies the trained machine learning model, and generates a CTR prediction.

Data Flows:

- User and ad data flow into the CTR Prediction System.
- The system outputs Predicted CTR results (click or no-click with probability) to the user through the web application.

Summary: The Level 0 DFD showcases the system as a single black box, focusing on input and output data flows, emphasizing system boundaries and user interaction.




Fig 4.2.1 Level 0 DFD

4.2.2 Level 1 DFD

The Level 1 DFD expands the single process into detailed sub-processes to illustrate the internal workflow of the CTR Prediction System, reflecting the practical pipeline you implemented using your Kaggle notebook and Streamlit app.

External Entities:

- User Data Source: Provides historical data for model training.
- Ad Data Source: Provides ad feature data.
- User (via Streamlit App): Inputs demographic, behavior, ad, and time-based data for prediction.

Processes:

1.Data Collection:

- Historical data is loaded from CSV files in Kaggle for model training.
- User input data is collected via Streamlit form during prediction.

2.Data Preprocessing:

- Missing value handling.
- Label encoding (city, country, ad topic).
- Frequency encoding (city, country).
- Gender encoding.
- Time feature extraction (day, hour, month, weekday).

3.Feature Engineering:

- Generates additional derived features and ensures consistent columns matching the trained model's structure.

4.Model Training and Selection:

- Applies machine learning algorithms (Random Forest, XGBoost) on the preprocessed training data.
- Performs hyperparameter tuning and selects the best-performing model based on evaluation metrics.

5.Evaluation:

- Calculates accuracy, MAE, RMSE, and feature importance for interpretability.
- Visualizes and validates model performance.

6.Deployment:

- The trained model is serialized using Joblib/Pickle and integrated into the Streamlit app.
- During user input, the model is loaded, and predictions are made in real-time.

7.Prediction and Output Display:

- The processed user input is passed to the model, which predicts whether the user will click on the ad and provides the click probability.

- Results and actionable insights (excellent, moderate, or poor targeting) are displayed to the user in the Streamlit interface, with optional feature importance explanations for interpretability.

Data Flows:

- Data Sources → Data Collection: Historical CSV data and real-time user input are collected.
- Data Collection → Data Preprocessing: Raw data is cleaned and transformed.
- Preprocessing → Feature Engineering: Additional meaningful features are created.
- Feature Engineered Data → Model Training: Used to train and evaluate ML models.
- Trained Model → Deployment: Saved and integrated for prediction.
- User Input (via Streamlit) → Preprocessing → Feature Engineering → Model Inference: For real-time prediction.
- Predicted CTR → User: Output displayed via Streamlit app.

Summary:

The Level 1 DFD breaks down the CTR Prediction System into clear stages representing your end-to-end pipeline:

- Data ingestion (historical & live),
- Cleaning and preprocessing,
- Model training and evaluation,
- Deployment for real-time prediction,
- User-friendly display of actionable results.

This structured flow ensures your system remains modular, scalable, and maintainable while accurately predicting CTR for real-time digital marketing optimization.

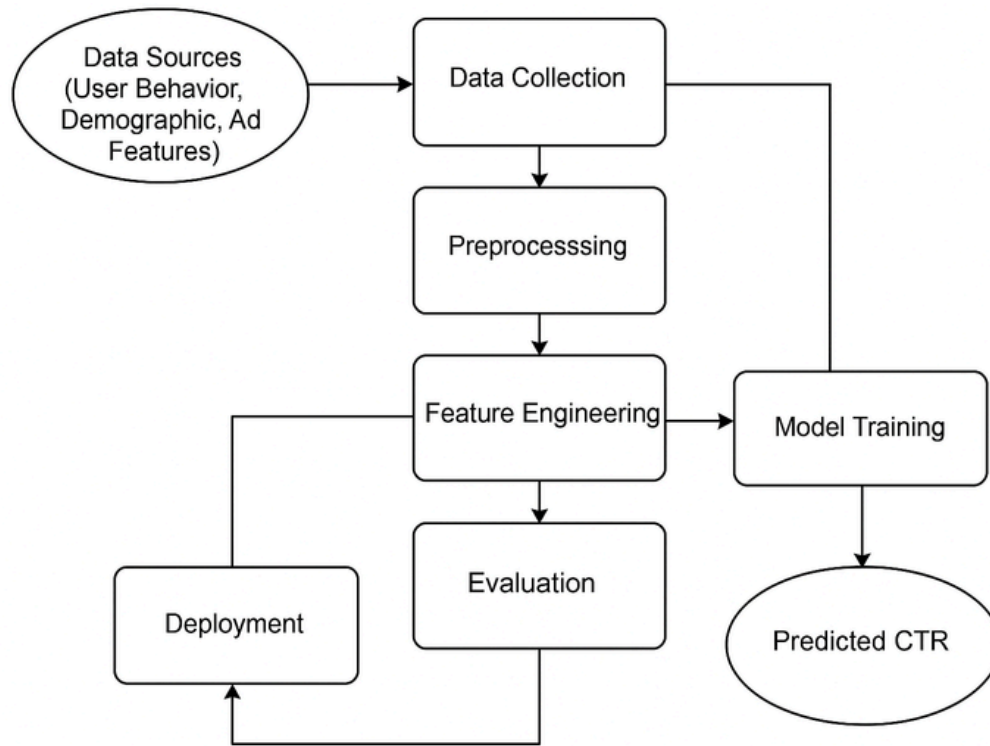


Fig 4.2.2 Level 1 DFD

4.3 SYSTEM ARCHITECTURE

The system architecture of the CTR Prediction Model defines a layered, modular structural framework that seamlessly connects all functional modules, ensuring smooth data flow from user input to real-time prediction while maintaining scalability and maintainability. It organizes the project into logical stages that mirror the practical machine learning pipeline implemented in the Kaggle notebook and deployed through the Streamlit application.

At the core, the architecture begins with the data ingestion layer, which collects historical data from CSV datasets (user demographics, user behavior, and ad attributes) during the model training phase and captures real-time user input data via an interactive Streamlit interface during prediction. This is followed by the data preprocessing layer, where the system systematically cleans the data, handles missing values, and applies transformations such as label encoding for categorical variables (city, country, ad topic), frequency encoding for location features, and the generation of derived time-based features (hour, day, month) to enhance predictive power.

System Architecture

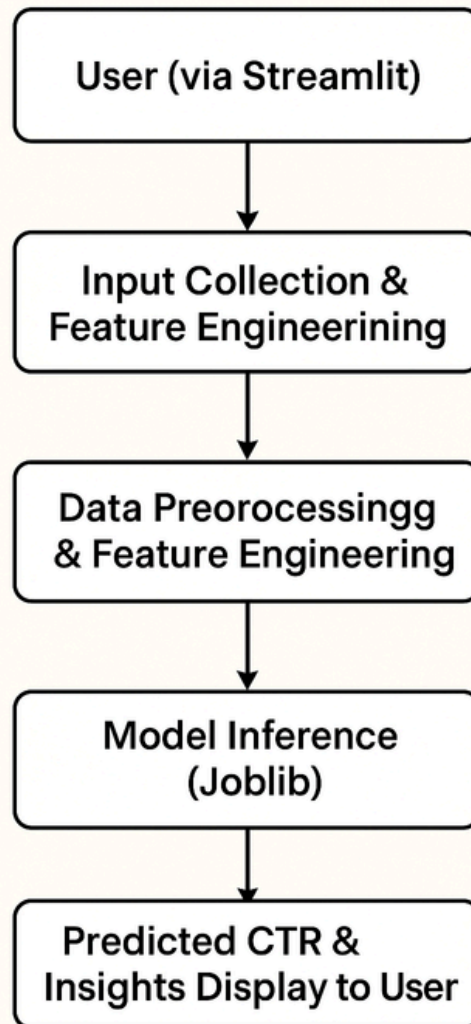


Fig 4.3 Architecture Diagram

Next, the feature engineering and selection layer ensures that the input features are structured consistently with the training pipeline, retaining only the most relevant columns and preparing the data in the exact order required by the trained model. This is followed by the model training and evaluation layer, where various machine learning algorithms, including Random Forest and XGBoost, are trained, tuned, and evaluated using metrics such as accuracy, MAE, and RMSE to identify the best-performing model, ensuring robust predictive capabilities.

Once the optimal model is selected, the model deployment layer uses Joblib to serialize and store the model along with preprocessing artifacts, enabling efficient reuse during real-time predictions. In the prediction phase, the Streamlit application loads the serialized model and artifacts, allowing the model inference layer to preprocess new user inputs consistently and generate CTR predictions instantly.

Finally, the application interface layer delivers results to the end user through a clean, intuitive Streamlit web interface. Here, users receive not only the CTR prediction (click/no-click) and its probability but also actionable insights into the predicted engagement potential and feature importance for transparency and learning. This architecture ensures that each stage in the pipeline operates systematically, reduces redundancy, and allows for easy debugging and future enhancements. Its modular design facilitates independent updates of data processing, modeling, and deployment components without impacting the rest of the system, ensuring a scalable, reusable, and production-ready CTR Prediction Model capable of practical deployment in real-world advertising environments.

4.3.1 Architecture Overview

The CTR Prediction Model architecture consists of four structured layers that ensure a smooth, modular, and scalable flow from data collection to real-time CTR prediction and actionable insights display. Each layer is designed to handle specific responsibilities, ensuring systematic execution aligned with your practical implementation.

1.Data Layer

This layer is responsible for ingesting all required data for both training and prediction:

Historical Data Collection:

- Loads ad click datasets from Kaggle in CSV format containing user demographics (age, gender, area income), user behavior (daily time spent on site, daily internet usage), ad-specific attributes (city, country, ad topic), and time-related details (hour, day, month).

- Utilizes pandas for structured data loading and initial exploration.

Real-Time User Input Collection:

- Uses Streamlit's form and input widgets (number_input, selectbox, slider) to collect user-provided data in real time, including:
 - Demographics: Age, gender, income.
 - Behavior: Daily site time, internet usage.
 - Ad context: City, country, ad topic.
 - Time: Hour of day, day of month, day of week, month.
- This allows dynamic data ingestion without requiring the user to handle CSV files.

2.Processing Layer

This layer is critical for transforming raw data into a structured, model-ready format, ensuring consistency between the training pipeline and real-time predictions.

Data Cleaning:

- Handles missing values systematically, ensuring no null values disrupt model training or predictions.

Categorical Feature Encoding:

- Label Encoding: City, country, and ad topic are encoded using label encoders created during training and saved within the model artifacts.
- Frequency Encoding: Additionally, city and country frequencies are computed from the training data and used to encode real-time inputs, capturing location-based significance for CTR prediction.
- Gender Encoding: Converts “Male” and “Female” into binary values (1 and 0).

Numerical Feature Processing:

- Features like daily site time, internet usage, and income are scaled and structured to match the trained model's input expectations.

Time Feature Engineering:

- Extracts and processes time-based features such as day of month, day of week, hour of day, and month, which help capture user activity patterns influencing ad click behavior.

Feature Alignment:

- Ensures real-time input data columns are reindexed and aligned precisely with the `feature_columns` saved during training, preventing column mismatch errors during model inference.

3.Model Layer

This layer is responsible for model training, evaluation, saving, and loading, ensuring accurate CTR prediction aligned with your practical workflow.

Model Training:

- Multiple machine learning algorithms were trained on the preprocessed Kaggle dataset:
 - Logistic Regression: Served as a baseline due to its interpretability and simplicity.
 - Decision Tree: Allowed capturing non-linear relationships between user behavior, demographics, and ad features.
 - Random Forest: Provided improved predictive performance through ensemble learning and reduced overfitting.
 - XGBoost: A gradient boosting algorithm capable of handling complex, non-linear patterns with high accuracy.
- All models were evaluated, and XGBoost demonstrated the highest accuracy and balanced performance across evaluation metrics, making it the final choice for deployment.

Model Evaluation:

- Each model was evaluated using:
 - Accuracy: To measure overall correct predictions.
 - Mean Absolute Error (MAE) and Root Mean Square Error (RMSE): To assess prediction error magnitudes.
- XGBoost outperformed other models in providing robust, consistent CTR predictions with the lowest error rates.

Model Saving:

- The final trained XGBoost model was serialized using Joblib along with:
 - Label encoders for categorical feature transformation.
 - Frequency maps for location-based features.

- The ordered list of feature_columns for consistent data alignment.
- Feature importance data for interpretability during deployment.

Model Loading:

- During deployment in the Streamlit application, the saved XGBoost model and preprocessing artifacts are loaded efficiently to allow instant predictions on new user input without the need for retraining, ensuring scalability and rapid response in real-time usage.

4. Application Layer

This layer implements a user-facing interface for real-time CTR prediction using Streamlit.

Interactive Web Application:

- Displays a clean, structured interface where users input data via forms and sliders, simplifying user interaction with the system.

Real-Time Inference:

- Preprocesses user inputs using the same transformations as in the training pipeline.
- Passes structured data into the loaded model to predict:
 - CTR (Click/No Click) as binary output.
 - Click Probability for detailed insight.

Actionable Insights and Interpretability:

- Displays the predicted probability with guidance:
 - High probability: Excellent targeting potential.
 - Moderate probability: Room for optimization.
 - Low probability: Poor targeting; ad refinement suggested.
- Presents feature importance dynamically to explain what influences the prediction, enhancing transparency for marketing teams.

User Feedback and Alerts:

- Uses Streamlit's success, warning, and error alerts to guide the user after prediction.

4.4 LIBRARIES AND TOOLS

The CTR Prediction Model leverages a set of carefully chosen libraries and tools to ensure efficient data processing, model training, evaluation, deployment, and real-time user interaction through a clean interface.

1. pandas

Used for data loading, manipulation, and analysis. It helps in:

- Reading CSV files from Kaggle datasets for model training.
- Structuring user input data during real-time predictions.
- Handling missing values and generating derived features during preprocessing.

2. scikit-learn

The core machine learning library used for:

- Model implementation: Logistic Regression, Decision Tree, Random Forest.
- Preprocessing: Label encoding for categorical variables, scaling features.
- Model evaluation: Computing accuracy, MAE, and RMSE.

3. matplotlib and seaborn

Used for data visualization and exploratory data analysis (EDA):

- Generating graphs for understanding feature distributions, correlations, and outlier detection.
- Visualizing model performance, feature importance, and results for interpretation during training and evaluation.

4. Logistic Regression

A linear classification algorithm that models the probability of a binary outcome (click or no-click) using a logistic (sigmoid) function. It assumes a linear relationship between the input features and the log-odds of the output, making it simple and fast for baseline predictions.

5. Decision Tree Classifier

A tree-based supervised learning algorithm that splits the dataset into subsets based on feature values using if-else rules. It builds a tree structure where each node represents a decision based on a feature, making decisions interpretable and capable of capturing non-linear relationships in the data.

6. Random Forest Classifier

An ensemble learning method that builds multiple Decision Trees during training and aggregates their outputs to improve accuracy and reduce overfitting. It uses bagging (bootstrap aggregation) to create diverse trees and takes the majority vote for the final prediction.

7. XGBoost Classifier

A gradient boosting algorithm that builds an ensemble of weak learners (typically decision trees) sequentially, where each new tree corrects the errors of the previous ones. It uses boosting with gradient descent optimization and incorporates regularization, making it efficient, scalable, and highly accurate for structured data prediction tasks.

Outcome: XGBoost demonstrated the highest accuracy during evaluation, making it the final choice for deployment in your Streamlit application.

8. joblib

Used for model serialization and deserialization:

- Saves the trained XGBoost model, label encoders, frequency maps, and feature columns efficiently as a .pkl file.
- Allows loading the saved model in the Streamlit app without retraining, ensuring fast real-time predictions.

9. Streamlit

Used to build the interactive web application for real-time CTR predictions:

- Provides an intuitive interface for users to input demographic, behavioral, and ad-related data.
- Displays CTR predictions, click probabilities, and actionable insights dynamically.
- Supports fast, local deployment for testing and sharing the project.

10. Kaggle Notebook

Used for model development, experimentation, and evaluation. Kaggle notebooks provided:

- A cloud-based, GPU/CPU-enabled environment for training models efficiently.
- Easy integration with datasets for preprocessing and feature engineering.
- Interactive execution and visualization for clear step-by-step experimentation.

11. VS Code (Visual Studio Code)

- Used to run and test the Streamlit app locally on your device before deployment. VS Code provided:
- A clean, organized development environment for writing and managing your Streamlit app code.
- Integrated terminal for running streamlit app locally for live testing.
- Ease of debugging and updating the user interface dynamically during development.

12. datetime

Used to handle time-related features:

- Extracts hour, day, and month features for inclusion in the dataset.
- Helps capture temporal patterns in user behavior affecting ad click probabilities.

4.5 MODULES

The CTR Prediction Model system is organized into clear, modular components to ensure systematic execution, maintainability, and scalability. Each module reflects a specific stage in your pipeline, from Kaggle-based model development to Streamlit deployment for real-time predictions.

1.Data Collection Module

Purpose:

Collects and organizes historical ad and user data for model training and captures real-time user inputs for predictions.

Implementation:

- In Kaggle Notebook:
 - Loads datasets containing user demographics (age, gender, area income), user behavior (daily time spent on site, daily internet usage), ad-related features (city, country, ad topic), and timestamps.
- In Streamlit App:
 - Uses Streamlit form widgets (number_input, selectbox, slider) to collect user input dynamically for real-time prediction.

Outcome:

Ensures data is ingested in a structured manner for consistent preprocessing and feature alignment.

2. Data Preprocessing Module

Purpose:

Prepares raw data for effective model training and prediction, ensuring consistency and quality.

Implementation:

- Handling Missing Values: Ensures no null values disrupt the pipeline.
- Label Encoding: Uses LabelEncoder for categorical variables (city, country, ad topic) based on encoders fitted during training.
- Frequency Encoding: Applies frequency mapping for city and country to capture regional significance.

- Gender Encoding: Converts gender into binary values (Male = 1, Female = 0).
- Time Feature Extraction: Derives hour, day, month, and weekday from timestamp data to capture temporal user behavior patterns.

Outcome:

Transforms raw data into a model-ready format, maintaining alignment with the trained model's expected feature structure.

3. Exploratory Data Analysis (EDA) Module

Purpose:

To understand data distribution, detect anomalies, and identify feature relationships before preprocessing and model building.

Implementation:

- Descriptive Statistics: Checked dataset size, null values, and summary statistics.
- Visualization: Used matplotlib and seaborn for:
 - Distribution plots of numerical features (age, income, site time, internet usage).
 - Count plots for categorical features (city, country, ad topic).
 - Correlation heatmaps to identify feature relationships.
 - Outlier detection using boxplots.

Outcome:

- Gained insights into user behavior patterns and ad feature distributions.
- Identified features needing transformation or encoding.
- Detected and handled outliers to improve model stability.

4. Feature Engineering Module

Purpose:

Enhances predictive performance by creating and structuring meaningful features.

Implementation:

- Creates new derived features (e.g., temporal features, frequency-encoded location features).

- Aligns user input columns to match the feature_columns used during training for consistent inference in the Streamlit app.
- Handles any missing or extra columns during real-time prediction to prevent errors.

Outcome:

Provides a robust, consistent feature set to improve the predictive accuracy of the model.

5. Model Training and Evaluation Module

Purpose:

Builds and evaluates machine learning models to identify the best-performing model for CTR prediction.

Implementation:

- Trains multiple models on the preprocessed data:
 - Logistic Regression
 - Decision Tree
 - Random Forest
 - XGBoost (Selected)
- Evaluates models using accuracy, MAE, and RMSE.
- Extracts and saves feature importance for interpretability.

Outcome:

Selects the XGBoost model due to its superior accuracy, preparing it for deployment.

6. Model Saving and Loading Module

Purpose:

Enables reusability of the trained model and preprocessing artifacts for real-time predictions.

Implementation:

- Uses Joblib to serialize:
 - The trained XGBoost model.
 - Label encoders for categorical features.

- Frequency maps for location encoding.
- Feature importance data.
- List of feature columns (feature_columns).
- Loads these artifacts in the Streamlit app for consistent inference without retraining.

Outcome:

Facilitates seamless integration of the model into the prediction pipeline, reducing computational overhead during deployment.

7. Prediction and Inference Module

Purpose:

Generates real-time CTR predictions based on user input data in the deployed environment.

Implementation:

- Preprocesses real-time user inputs to align with the trained model pipeline.
- Uses the loaded XGBoost model to predict:
 - CTR (Click / No Click) as binary output.
 - Click probability for detailed insight.

Outcome:

Delivers immediate, accurate CTR predictions to users, aiding targeted ad decisions.

8. Visualization and User Interface Module

Purpose:

Provides a user-friendly, interactive interface for users to input data and receive actionable predictions.

Implementation:

- Built using Streamlit, providing:
 - Forms for user input collection.
 - Display of predicted CTR status (Will Click / Won't Click).
 - Display of click probability.
 - Visualization of feature importance for interpretability.
 - Actionable insights (Excellent, Moderate, Poor targeting suggestions).

Outcome:

Enhances the user experience, allowing non-technical users to utilize the CTR prediction system effectively.

4.6 PICKLE/JOBLIB MODULE

The Pickle/Joblib Module in the CTR Prediction Model system is crucial for saving and reusing the trained machine learning model and preprocessing artifacts without retraining, enabling real-time, scalable deployment in the Streamlit application.

What are Pickle and Joblib?

- **Pickle:** A Python module used for serializing and deserializing Python objects, allowing models and objects to be saved in a .pkl file and reloaded later.
- **Joblib:** An optimized alternative to Pickle, specifically designed for efficiently serializing large numpy arrays and machine learning models, reducing file size and loading time.

Why Joblib was used in this project?

Convenience: Simplifies saving and loading trained models and preprocessing encoders seamlessly.

Integration: Easily integrates with scikit-learn and XGBoost, ensuring compatibility and stability during deployment.

Implementation in the CTR Prediction Model

During model training in the Kaggle Notebook:

- After identifying XGBoost as the best-performing model, it, along with essential preprocessing artifacts, was saved using Joblib:
 - Trained XGBoost model.
 - Label Encoders for city, country, and ad topic.
 - Frequency Maps for location features.
 - Feature Importance Data for interpretability.
 - Ordered list of feature_columns to maintain consistent structure during inference.

During real-time prediction in the Streamlit app:

- The saved .pkl file is loaded at the start of the application:
- User inputs are preprocessed using the loaded encoders and mappings to match the training pipeline.
- The preprocessed input is passed to the loaded XGBoost model to generate CTR predictions instantly without retraining.

Advantages in Your Project

- **Enables Real-Time Predictions:** No need to retrain models for every new prediction.
- **Consistency:** Ensures the same encoders and features used during training are applied during deployment, avoiding mismatch errors.
- **Scalable Deployment:** Facilitates deploying the CTR Prediction Model within the Streamlit application efficiently.
- **Saves Time and Resources:** Eliminates redundant computation and leverages the trained model directly for user-facing predictions.

CHAPTER 5

CONCLUSION

5.1 FUTURE SCOPE

While the CTR Prediction Model using Machine Learning has achieved its goal of accurately predicting ad click-through rates and aiding advertisers in optimizing campaigns, there are several areas for future enhancement and expansion:

1.Integration of Deep Learning Models:

Exploring deep neural networks (DNN) or hybrid models like DeepFM could capture complex, high-order feature interactions, potentially improving prediction accuracy, especially on large-scale datasets.

2. Real-Time Data Streaming:

Currently, the system uses batch data for training and manual input for prediction. Future development can include real-time data streaming using Kafka or similar pipelines, allowing continuous model updates and predictions as new user and ad data become available.

3. Automated Model Retraining:

User behavior and ad effectiveness change over time. Incorporating scheduled or triggered model retraining pipelines would help maintain accuracy without manual intervention, ensuring the model adapts to changing data trends.

4. A/B Testing in Live Environments:

Deploying the model within an actual ad-serving system and conducting A/B testing will allow practical evaluation of its impact on ad click rates, engagement, and ROI.

5. Advanced Explainability Tools:

Integrating SHAP or LIME explainability frameworks can provide detailed feature-level explanations for each prediction, enhancing advertiser trust and interpretability of model decisions.

6. API Deployment:

Packaging the model as a REST API using FastAPI or Flask would enable seamless integration with ad platforms, mobile apps, or marketing dashboards for scalable real-world deployment.

7. Expanding Feature Sets:

Including additional features such as user device type, previous ad engagement, session time, and contextual signals (weather, events, location granularity) could improve prediction relevance and precision.

8. Multi-Platform CTR Prediction:

Extending the system to handle cross-platform CTR prediction (web, mobile, social media channels) would broaden its applicability in diverse advertising ecosystems.

9. UI Enhancements:

The Streamlit app can be improved by:

- Adding data validation for input fields.
- Visualizing user-specific feature importance.
- Providing suggestions for ad improvement based on predicted CTR.

5.2 CONCLUSION

The CTR Prediction Model using Machine Learning effectively demonstrates how data-driven approaches can enhance the efficiency and accuracy of digital advertising. By leveraging user demographics, behavioral data, and advertisement features, the system predicts the likelihood of a user clicking on an advertisement, thereby assisting advertisers in optimizing their campaigns and maximizing return on investment.

Throughout this project, a systematic approach was followed, beginning with data collection and exploratory data analysis (EDA) to understand data

distribution and patterns. Preprocessing and feature engineering ensured the data was clean and structured for machine learning models. Various models, including Logistic Regression, Decision Tree, Random Forest, and XGBoost, were trained and evaluated, with XGBoost ultimately selected due to its superior predictive performance and accuracy.

The trained model and preprocessing pipeline were seamlessly integrated into a Streamlit web application, allowing for real-time CTR predictions through a user-friendly interface. The use of Joblib ensured scalable and efficient model deployment, enabling instant predictions without retraining. Users can input new data to receive immediate feedback on the likelihood of an ad being clicked, along with actionable insights and interpretability through feature importance visualization.

This project showcases how machine learning can be practically applied to solve real-world problems in digital marketing, enhancing ad targeting, reducing unnecessary expenditure, and improving user engagement by delivering relevant advertisements. It also emphasizes the importance of clean data pipelines, modular system architecture, and user-friendly deployment in building scalable, maintainable data science projects.

In conclusion, the CTR Prediction Model bridges the gap between theoretical learning and practical application, providing a robust, efficient, and accessible solution for predictive ad targeting in the digital marketing ecosystem.

CHAPTER 6

APPENDICES

6.1 OUTPUT

The screenshot shows the top part of a web application titled "Ad Click-Through Rate (CTR) Prediction". Below the title is a subtitle "Predict whether a user will click on an advertisement". The interface is divided into three main sections: "User Demographics", "User Behavior", and "Location & Content". The "User Demographics" section includes input fields for "Age" (set to 60), "Gender" (set to Male), and "Area Income (\$)" (set to 30000.00). The "User Behavior" section includes input fields for "Daily Time Spent on Site (minutes)" (set to 20.00) and "Daily Internet Usage (minutes)" (set to 120.00). The "Location & Content" section is partially visible at the bottom. On the right side of the interface, there is a "Deploy" button and a vertical scrollbar.

Fig 6.1.1 Data Input 1

The screenshot shows the "Location & Content" section of the web application. It includes input fields for "City" (set to Andrewborough), "Country" (set to Afghanistan), and "Ad Topic" (set to Advanced 24/7 productivity). There are also sliders for "Hour of Day" (set to 0), "Day of Month" (set to 19), "Day of Week (0=Monday)" (set to 6), and "Month" (set to 3). A "Predict CTR" button is located at the bottom left of this section. On the right side of the interface, there is a "Deploy" button and a vertical scrollbar.

Fig 6.1.2 Data Input 2

6.1 OUTPUT

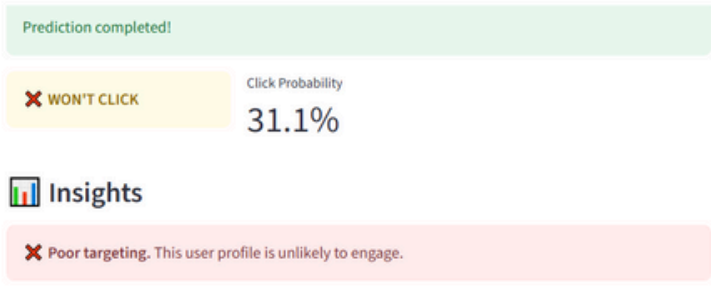


Fig 6.1.3 Poor Targeting

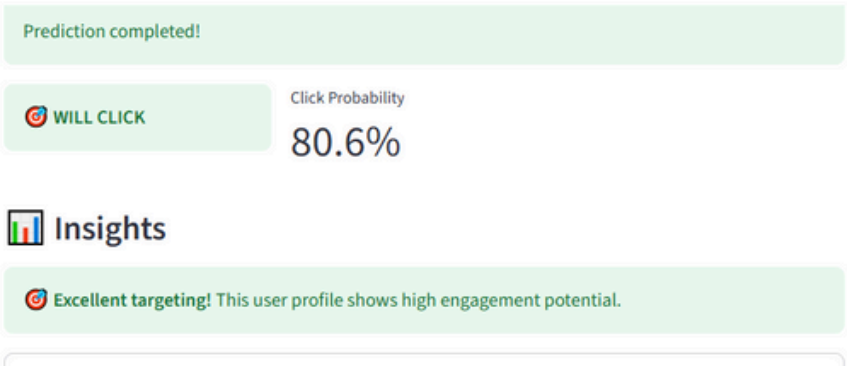


Fig 6.1.4 Excellent Targeting

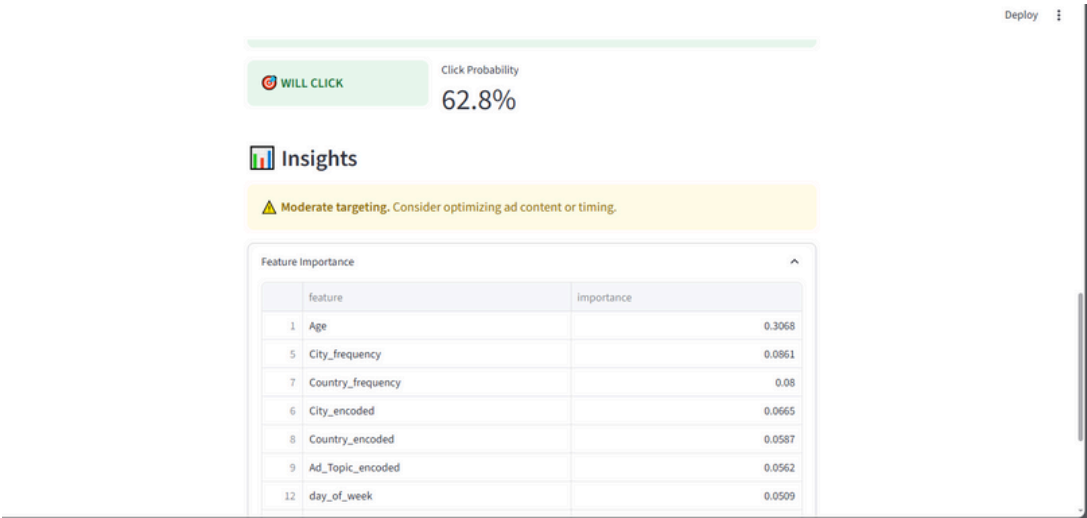


Fig 6.1.5 Moderate Targeting

CHAPTER 7

REFERENCES

- [1] H. He, D. McAuley and J. Leskovec, "Ups and Downs: Modeling the Visual Evolution of Fashion Trends with One-Class Collaborative Filtering," Proceedings of the 23rd International Conference on World Wide Web, Seoul, Korea, 2014, pp. 507-517.
- [2] A. Gupta, P. Jain, and N. Varma, "Practical Lessons from Predicting Clicks on Ads at Facebook," Proceedings of the Eighth ACM International Conference on Web Search and Data Mining (WSDM), Shanghai, China, 2015, pp. 455-464.
- [3] M. Richardson, E. Dominowska and R. Ragno, "Predicting Clicks: Estimating the Click-Through Rate for New Ads," Proceedings of the 16th International Conference on World Wide Web (WWW), Banff, Canada, 2007, pp. 521-530.
- [4] W. Zhang, T. Du, and J. Wang, "Deep Learning over Multi-Field Categorical Data: A Case Study on User Response Prediction," arXiv preprint arXiv:1601.02376, 2016.
- [5] Y. Juan, Y. Zhuang, W. Chin and C. Lin, "Field-aware Factorization Machines for CTR Prediction," Proceedings of the 10th ACM Conference on Recommender Systems (RecSys), Boston, MA, USA, 2016, pp. 43-50.
- [6] K. Gai, B. Zhu, Y. Liu, L. Zhu, and M. Qiu, "Privacy-Preserving Multi-Channel Communication in Edge-of-Things," Future Generation Computer Systems, vol. 85, pp. 190-200, 2018.

- [7] H. Guo, R. Tang, Y. Ye, Z. Li and X. He, "DeepFM: A Factorization-Machine based Neural Network for CTR Prediction," Proceedings of the 26th International Joint Conference on Artificial Intelligence (IJCAI), Melbourne, Australia, 2017, pp. 1725-1731.
- [8] H. B. McMahan et al., "Ad Click Prediction: a View from the Trenches," Proceedings of the 19th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD), Chicago, IL, USA, 2013, pp. 1222-1230.
- [9] S. Rendle, "Factorization Machines with libFM," ACM Transactions on Intelligent Systems and Technology (TIST), vol. 3, no. 3, pp. 57:1-57:22, 2012.
- [10] X. He, T. Chen, and T. Jin, "Practical Lessons from Predicting Clicks on Ads at Facebook," Proceedings of the 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD), New York, USA, 2014, pp. 1039-1048.